


Training Curves: Loss and Accuracy (Phase 1: frozen backbone; Phase 2: fine-tuned)



Solid = train, Dashed = validation. Dotted line separates Phase 1 from Phase 2.
DenseNet121 uses focal loss ($\gamma = 2.0$, class-weighted); loss values not comparable across models.